ARM Processor Architecture

Some Slides are Adopted from NCTU
IP Core Design
Some Slides are Adopted from NTU
Digital SIP Design Project

Outline

❑ARM Core Family

❑ARM Processor Core

❑Introduction to Several ARM processors

❑Memory Hierarchy

❑Software Development

❑Summary

ARM Core Family

ARM Core Family

Application Cores	Embedded Cores	Secure Cores
ARM Cortex-A8	ARM Cortex-M3	SecurCore SC100
ARM1020E	ARM1026EJ-S	SecurCore SC110
ARM1022E	ARM1156T2(F)-S	SecurCore SC200
ARM1026EJ-S	ARM7EJ-S	SecurCore SC210
ARM11 MPCore	ARM7TDMI	
ARM1136J(F)-S	ARM7TDMI-S	
ARM1176JZ(F)-S	ARM946E-S	
ARM720T	ARM966E-S	
ARM920T	ARM968E-S	
ARM922T	ARM996HS	
ARM926EJ-S		

Product Code Demystified

- ❑ T: Thumb
- ❑ D: On-chip debug support
- ❑ M: Enhanced multiplier
- ❑ I: Embedded ICE hardware
- ❑ T2: Thumb-2
- ❑ S: Synthesizable code
- ❑ E: Enhanced DSP instruction set
- ❑ J: JAVA support, Jazelle
- ❑ Z: Should be TrustZone?
- ❑ F: Floating point unit
- ❑ H: Handshake, clockless design for synchronous or asynchronous design

ARM Processor Cores (1/4)

❑ ARM processor core + cache + MMU

→ ARM CPU cores

❑ ARM6 → ARM7

– 3-stage pipeline

– Keep its instructions and data in the same memory system

– **T**humb 16-bit compressed instruction set

– On-chip **D**ebug support, enabling the processor to halt in response to a debug request

– Enhanced **M**ultiplier, 64-bit result

– Embedded **I**CE hardware, give on-chip breakpoint and watchpoint support

ARM Processor Cores (2/4)

❑ ARM8 → ARM9

      → ARM10

❑ ARM9

- 5-stage pipeline (130 MHz or 200MHz)
- Using separate instruction and data memory ports

❑ ARM 10 (1998. Oct.)

- High performance, 300 MHz
- Multimedia digital consumer applications
- Optional vector floating-point unit

ARM Processor Cores (3/4)

❑ ARM11 (2002 Q4)

- 8-stage pipeline
- Addresses a broad range of applications in the wireless, consumer, networking and automotive segments
- Support media accelerating extension instructions
- Can achieve 1GHz
- Support AXI

❑ SecurCore Family

– Smart card and secure IC development

ARM Processor Cores (4/4)

❑ Cortex Family

- – Provides a large range of solutions optimized around specific market applications across the full performance spectrum
- – ARM Cortex-A Series, applications processors for complex OS and user applications.
    - • Supports the ARM, Thumb and Thumb-2 instruction sets
- – ARM Cortex-R Series, embedded processors for real-time systems.
    - • Supports the ARM, Thumb, and Thumb-2 instruction sets
- – ARM Cortex-M Series, deeply embedded processors optimized for cost sensitive applications.
    - • Supports the Thumb-2 instruction set only

ARM Processor Core

❏ Version 1

– The first ARM processor, developed at Acorn Computers Limited 1983-1985

– 26-bit address, no multiply or coprocessor support

❏ Version 2

– Sold in volume in the Acorn Archimedes and A3000 products

– 26-bit addressing, including 32-bit result multiply and coprocessor

❏ Version 2a

– Coprocessor 15 as the system control coprocessor to manage cache

– Add the atomic load store (SWP) instruction

ARM Architecture Version (2/6)

❑ Version 3

- First ARM processor designed by ARM Limited (1990)
- ARM6 (macro cell)

  ARM60 (stand-alone processor)

  ARM600 (an integrated CPU with on-chip cache, MMU, write buffer)

  ARM610 (used in Apple Newton)
- 32-bit addressing, separate CPSR and SPSRs
- Add the undefined and abort modes to allow coprocessor emulation and virtual memory support in supervisor mode

❑ Version 3M

- Introduce the signed and unsigned multiply and multiply-accumulate instructions that generate the full 64-bit result

ARM Architecture Version (3/6)

❑ Version 4
- – Add the signed, unsigned half-word and signed byte load and store instructions
- – Reserve some of SWI space for architecturally defined operation
- – System mode is introduced

❑ Version 4T
- – 16-bit Thumb compressed form of the instruction set is introduced

❑ Version 5T
- – Introduced recently, a superset of version 4T adding the BLX, CLZ and BRK instructions

❑ Version 5TE
- – Add the signal processing instruction set extension
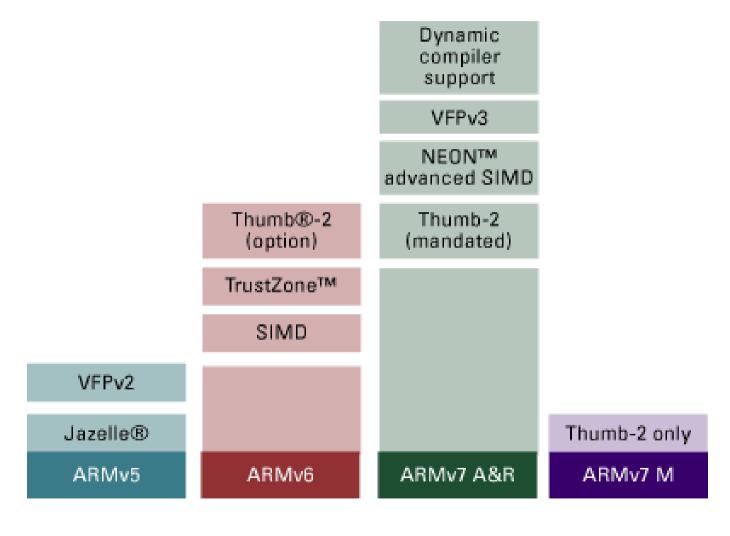
ARM Architecture Version (4/6)

❑ Version 6

– **Media processing extensions (SIMD)**
- **2x faster MPEG4 encode/decode**
- **2x faster audio DSP**

– Improved cache architecture
- Physically addressed caches
- Reduction in cache flush/refill
- Reduced overhead in context switches

– Improved exception and interrupt handling
- Important for improving performance in real-time tasks

– Unaligned and mixed-endian data support
- Simpler data sharing, application porting and saves memory
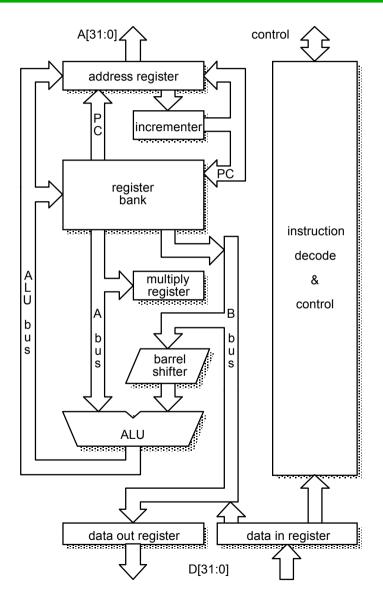
ARM Architecture Version (5/6)

ARM Architecture Version (6/6)

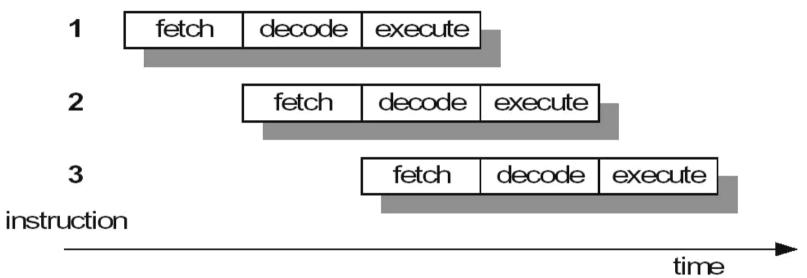| Core	Architecture
ARM1	v1
ARM2	v2
ARM2as, ARM3	v2a
ARM6, ARM600, ARM610	v3
ARM7, ARM700, ARM710	v3
ARM7TDMI, ARM710T, ARM720T, ARM740T	v4T
StrongARM, ARM8, ARM810	v4
ARM9TDMI, ARM920T, ARM940T	V4T
ARM9E-S, ARM10TDMI, ARM1020E	v5TE
ARM10TDMI, ARM1020E	v5TE
ARM11 MPCore, ARM1136J(F)-S, ARM1176JZ(F)-S	v6
Cortex-A/R/M	v7

3-Stage Pipeline ARM Organization

- **Register Bank**
  - 2 read ports, 1 write ports, access any register
  - 1 additional read port, 1 additional write port for r15 (PC)

- **Barrel Shifter**
  - Shift or rotate the operand by any number of bits

- **ALU**

- **Address register and incrementer**

- **Data Registers**
  - Hold data passing to and from memory

- **Instruction Decoder and Control**

3-Stage Pipeline (1/2)

❑ Fetch
- – The instruction is fetched from memory and placed in the instruction pipeline

❑ Decode
- – The instruction is decoded and the datapath control signals prepared for the next cycle

❑ Execute
- – The register bank is read, an operand shifted, the ALU result generated and written back into destination register
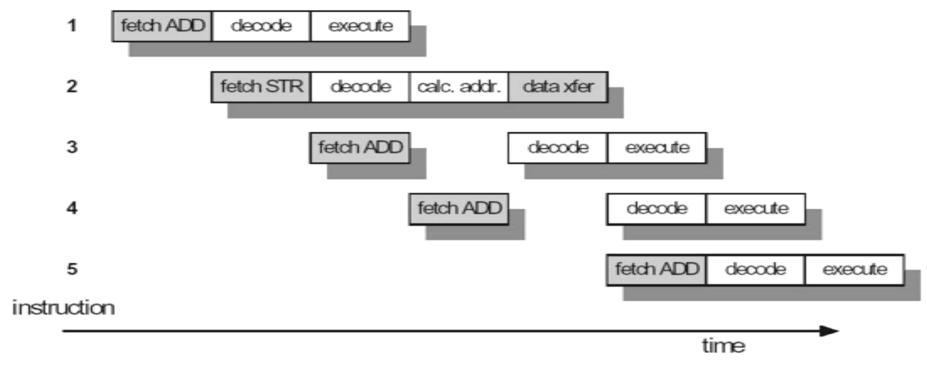
3-Stage Pipeline (2/2)

❏ At any time slice, 3 different instructions may occupy each of these stages, so the hardware in each stage has to be capable of independent operations

❏ When the processor is executing data processing instructions , the latency = **3** cycles and the throughput = **1** instruction/cycle

Multi-Cycle Instruction

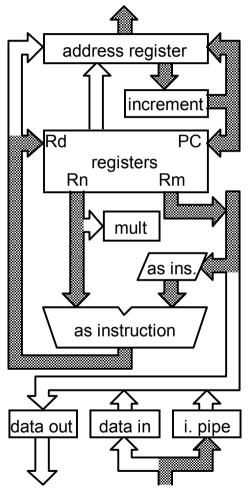- ❑ Memory access (fetch, data transfer) in every cycle
- ❑ Datapath used in every cycle (execute, address calculation, data transfer)
- ❑ Decode logic generates the control signals for the data path use in next cycle (decode, address calculation)

Data Processing Instruction

(a) register - register operations   *(b) register - immediate operations*
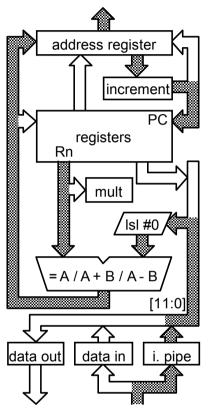
❑ All operations take place in a single clock cycle

Data Transfer Instructions

(a) 1st cycle - compute address   *(b) 2nd cycle - store data & auto-index*
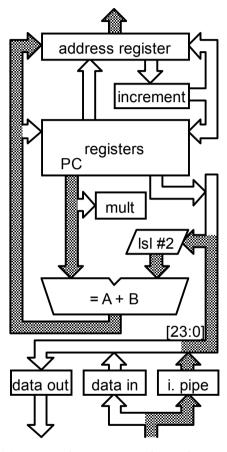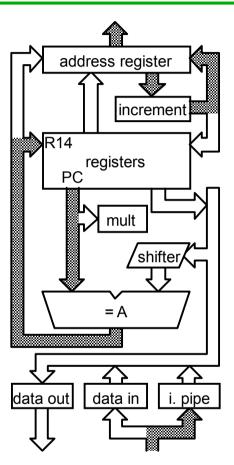
❑ Computes a memory address similar to a data processing instruction
❑ Load instruction follows a similar pattern except that the data from memory only gets as far as the 'data in' register on the 2nd cycle and a 3rd cycle is needed to transfer the data from there to the destination register

Branch Instructions

(a) 1st cycle - compute branch target

(b) 2nd cycle - save return address

❑ The third cycle, which is required to complete the pipeline refilling, is also used to mark the small correction to the value stored in the link register in order that is points directly at the instruction which follows the branch

Branch Pipeline Example

Cycle				1	2	3	4	5
address	opeation							
0x8000	BL	fetch	decode	execute	linkret	adjust		
0x8004	X		fetch	decode				
0x8008	XX			fetch				
0x8FEC	ADD				fetch	decode	execute	
0x8FF0	SUB					fetch	decode	execute
0x8FF4	MOV						fetch	decode
								fetch

❑ Breaking the pipeline

❑ Note that the core is executing in the ARM state
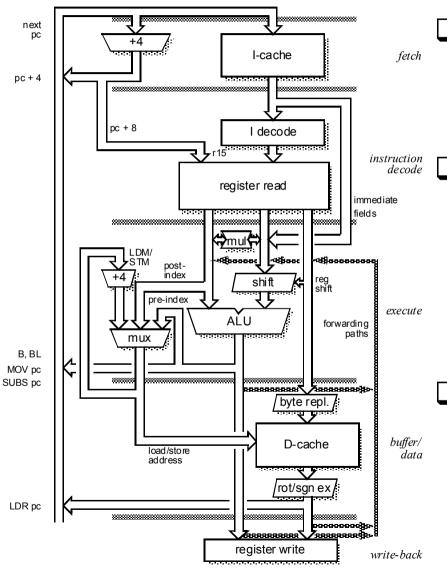
5-Stage Pipeline ARM Organization

❑ $T_{prog} = N_{inst} * CPI / f_{clk}$

– T_{prog}: the time that executes a given program

– N_{inst}: the number of ARM instructions executed in the program => compiler dependent

– CPI: average number of clock cycles per instructions => hazard causes pipeline stalls

– f_{clk}: frequency

❑ Separate instruction and data memories => **5** stage pipeline

❑ Used in ARM9TDMI

5-Stage Pipeline Organization (1/2)

fetch

instruction decode

execute

buffer/ data

write-back

Diagram labels: next pc, +4, I-cache, pc + 4, pc + 8, I decode, r15, register read, immediate fields, mul, LDM/STM, post-index, +4, shift, reg shift, pre-index, ALU, mux, B, BL, MOV pc, SUBS pc, forwarding paths, byte repl., load/store address, D-cache, rot/sgn ex, LDR pc, register write

❑ **Fetch**

– The instruction is fetched from memory and placed in the instruction pipeline

❑ **Decode**
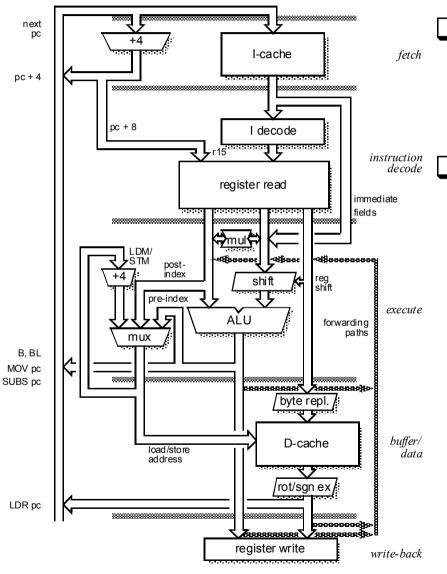
– The instruction is decoded and register operands read from the register files. There are 3 operand read ports in the register file so most ARM instructions can source all their operands in one cycle

❑ **Execute**

– An operand is shifted and the ALU result generated. If the instruction is a load or store, the memory address is computed in the ALU

5-Stage Pipeline Organization (2/2)

❑ Buffer/Data
  – Data memory is accessed if required. Otherwise the ALU result is simply buffered for one cycle

❑ Write back
  – The result generated by the instruction are written back to the register file, including any data loaded from memory

Pipeline Hazards

❑ There are situations, called *hazards*, that prevent the next instruction in the instruction stream from being executing during its designated clock cycle. Hazards reduce the performance from the ideal speedup gained by pipelining.

❑ There are three classes of hazards:

– **Structural Hazards**
  • They arise from resource conflicts when the hardware cannot support all possible combinations of instructions in simultaneous overlapped execution.

– **Data Hazards**
  • They arise when an instruction depends on the result of a previous instruction in a way that is exposed by the overlapping of instructions in the pipeline.

– **Control Hazards**
  • They arise from the pipelining of branches and other instructions that change the PC

Structural Hazards

❏ When a machine is pipelined, the overlapped execution of instructions requires pipelining of functional units and duplication of resources to allow all possible combinations of instructions in the pipeline.

❏ If some combination of instructions cannot be accommodated because of a *resource conflict*, the machine is said to have a **structural hazard**.

Example

❑ A machine has shared a **single-memory** pipeline for data and instructions. As a result, when an instruction contains a data-memory reference (load), it will conflict with the instruction reference for a later instruction (instr 3):

instr	Clock cycle number							
	1	2	3	4	5	6	7	8
load	IF	ID	EX	**MEM**	WB			
Instr 1		IF	ID	EX	MEM	WB		
Instr 2			IF	ID	EX	MEM	WB	
Instr 3				**IF**	ID	EX	MEM	WB

Solution (1/2)

❑ To resolve this, we *stall* the pipeline for one clock cycle when a data-memory access occurs. The effect of the stall is actually to occupy the resources for that instruction slot. The following table shows how the stalls are actually implemented.

	Clock cycle number								
instr	1	2	3	4	5	6	7	8	9
load	IF	ID	EX	**MEM**	WB				
Instr 1		IF	ID	EX	MEM	WB			
Instr 2			IF	ID	EX	MEM	WB		
Instr 3				stall	IF	ID	EX	MEM	WB

Solution (2/2)

❑ Another solution is to use separate instruction and data memories.

❑ ARM belongs to the **Harvard** architecture, so it does not suffer from this hazard

Data Hazards

❑ **Data hazards** occur when the pipeline changes the order of read/write accesses to operands so that the order differs from the order seen by sequentially executing instructions on the unpipelined machine.

		\multicolumn{9}{c}{Clock cycle number}								
		1	2	3	4	5	6	7	8	9
ADD	R1,R2,R3	IF	ID	EX	MEM	**WB**				
SUB	R4,R5,R1		IF	ID_{sub}	EX	MEM	WB			
AND	R6,R1,R7			IF	ID_{and}	EX	MEM	WB		
OR	R8,R1,R9				IF	ID_{or}	EX	MEM	WB	
XOR	R10,R1,R11					IF	ID_{xor}	EX	MEM	WB

Forwarding

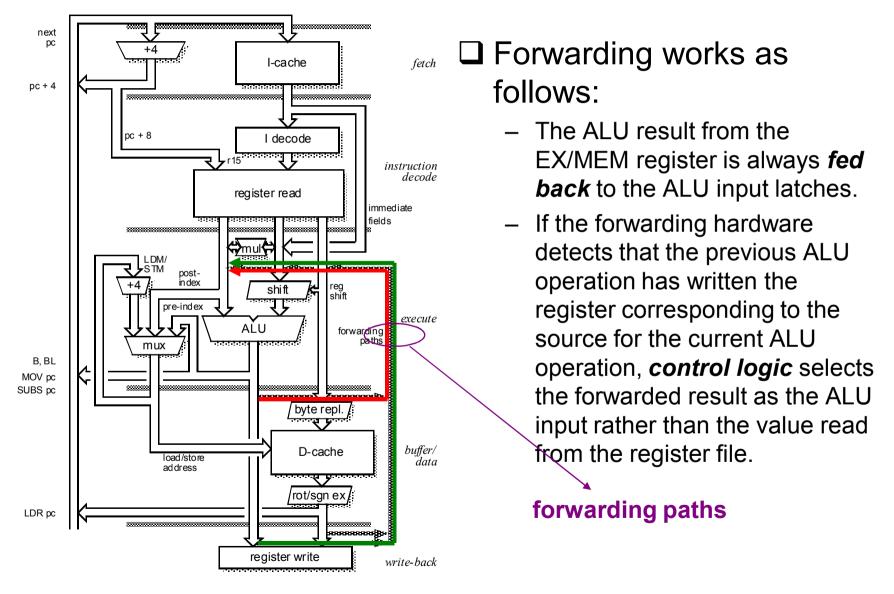❑ The problem with data hazards, introduced by this sequence of instructions can be solved with a simple hardware technique called *forwarding*.

		Clock cycle number						
		1	2	3	4	5	6	7
ADD	R1,R2,R3	IF	ID	EX	MEM	**WB**		
SUB	R4,R5,R1		IF	ID$_{sub}$	EX	MEM	WB	
AND	R6,R1,R7			IF	ID$_{and}$	EX	MEM	WB

Forwarding Architecture

- ❑ Forwarding works as follows:
  - The ALU result from the EX/MEM register is always *fed back* to the ALU input latches.
  - If the forwarding hardware detects that the previous ALU operation has written the register corresponding to the source for the current ALU operation, *control logic* selects the forwarded result as the ALU input rather than the value read from the register file.

  forwarding paths

Forward Data

Clock cycle number								
		1	2	3	4	5	6	7
ADD	R1,R2,R3	IF	ID	EX$_{add}$	MEM$_{add}$	WB		
SUB	R4,R5,R1		IF	ID	EX$_{sub}$	MEM	WB	
AND	R6,R1,R7			IF	ID	EX$_{and}$	MEM	WB

❑ The first forwarding is for value of **R1** from **EX$_{add}$** to **EX$_{sub}$**.
The second forwarding is also for value of **R1** from **MEM$_{add}$** to **EX$_{and}$**.
This code now can be executed without stalls.

❑ Forwarding can be generalized to include passing the result directly to the functional unit that requires it: a result is forwarded from the output of one unit to the input of another, rather than just from the result of a unit to the input of the same unit.

Without Forward

Clock cycle number										
		1	2	3	4	5	6	7	8	9
ADD	R1,R2,R3	IF	ID	EX	MEM	**WB**				
SUB	R4,R5,R1		IF	*stall*	*stall*	ID_{sub}	EX	MEM	WB	
AND	R6,R1,R7			*stall*	*stall*	IF	ID_{and}	EX	MEM	WB

Data Forwarding

❏ Data dependency arises when an instruction needs to use the result of one of its predecessors before the result has returned to the register file => pipeline hazards

❏ Forwarding paths allow results to be passed between stages as soon as they are available

❏ 5-stage pipeline requires each of the three source operands to be forwarded from any of the intermediate result registers

❏ Still one load stall

```
LDR rN, […]
ADD r2,r1,rN  ;use rN immediately
```

- One stall
- Compiler rescheduling

Stalls are Required

		1	2	3	4	5	6	7	8
LDR	R1,@(R2)	IF	ID	EX	MEM	WB			
SUB	R4,R1,R5		IF	ID	EX$_{sub}$	MEM	WB		
AND	R6,R1,R7			IF	ID	EX$_{and}$	MEM	WB	
OR	R8,R1,R9				IF	ID	EXE	MEM	WB

❑ The load instruction has a delay or latency that cannot be eliminated by forwarding alone.

The Pipeline with one Stall

		1	2	3	4	5	6	7	8	9
LDR	R1,@(R2)	IF	ID	EX	MEM	WB				
SUB	R4,R1,R5		IF	ID	stall	EX$_{sub}$	MEM	WB		
AND	R6,R1,R7			IF	stall	ID	EX	MEM	WB	
OR	R8,R1,R9				stall	IF	ID	EX	MEM	WB

❑ The only necessary forwarding is done for R1 from **MEM** to **EX$_{sub}$**.

LDR Interlock

Cycle			1	2	3	4	5	6	7	8	9
Operation											
ADD	R1, R1, R2	F	D	E		W					
SUB	R3, R4, R1		F	D	E		W				
LDR	R4, [R7]			F	D	E	M	W			
ORR	R8, R3, R4				F	D	I	E		W	
AND	R6, R3, R1					F	I	D	E		W
EOR	R3, R1, R2							F	D	E	W

F - Fetch   D - Decode   E - Excute   I - Interlock   M - Memory
W - Writeback

❑ In this example, it takes 7 clock cycles to execute 6 instructions, CPI of 1.2

❑ The LDR instruction immediately followed by a data operation using the same register cause an interlock

Optimal Pipelining

Cycle		1	2	3	4	5	6	7	8	9
Operation										
ADD	R1, R1, R2	F	D	E		W				
SUB	R3, R4, R1		F	D	E		W			
LDR	R4, [R7]			F	D	E	M	W		
AND	R6, R3, R1				F	D	E		W	
ORR	R8, R3, R4					F	D	E		W
EOR	R3, R1, R2						F	D	E		W

F - Fetch   D - Decode   E - Excute   I - Interlock   M - Memory

W - Writeback

❑ In this example, it takes 6 clock cycles to execute 6 instructions, CPI of 1

❑ The LDR instruction does not cause the pipeline to interlock

LDM Interlock (1/2)

Cycle		1	2	3	4	5	6	7	8	9	10
Operation											
LDMLA	R13!, {R0-R3}	F	D	E	M	MW	MW	MW	W		
SUB	R9, R7, R2		F	D	I	I	I	E		W	
STR	R4, [R9]			F	I	I	I	D	E	M	W
ORR	R8, R4, R3					F	D	E		W	
AND	R6, R3, R1						F	D	E		W

F - Fetch   D - Decode   E - Excute   I - Interlock   M - Memory

ME - Simultaneous Memory and Writeback   W - Writeback

❑ In this example, it takes 8 clock cycles to execute 5 instructions, CPI of 1.6

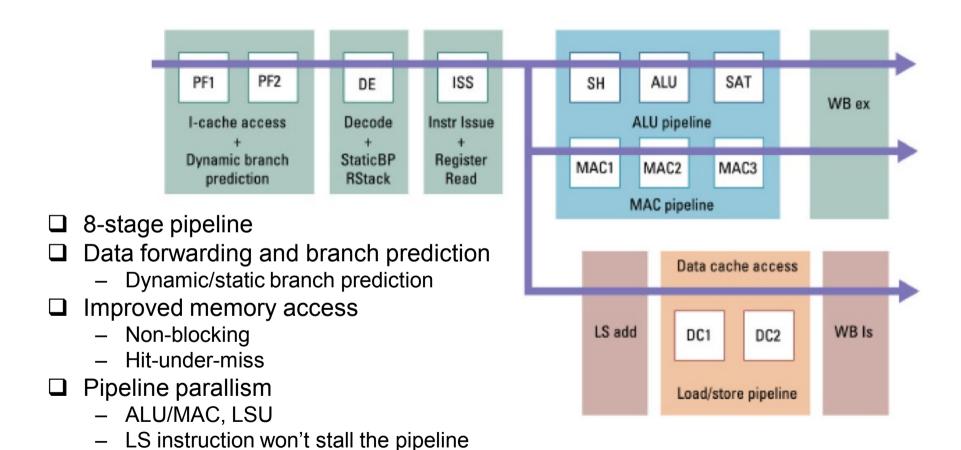❑ During the LDM there are parallel memory and writeback cycles

LDM Interlock (2/2)

Cycle		1	2	3	4	5	6	7	8	9	10
Operation											
LDMLA	R13!, {R0-R3}	F	D	E	M	MW	MW	MW	W		
SUB	R9, R7, R3		F	D	I	I	I	I	E		W
STR	R4, [R9]			F	I	I	I	I	D	E	M	W
ORR	R8, R4, R3						F	D	E		W
AND	R6, R3, R1							F	D	E	

F - Fetch   D - Decode   E - Excute   I - Interlock   M - Memory

ME - Simultaneous Memory and Writeback   W - Writeback

❑ In this example, it takes 9 clock cycles to execute 5 instructions, CPI of 1.8

❑ The SUB incurs a further cycle of interlock due to it using the highest specified register in the LDM instruction

8-Stage Pipeline (v6 Architecture)

- ❑ 8-stage pipeline
- ❑ Data forwarding and branch prediction
  - – Dynamic/static branch prediction
- ❑ Improved memory access
  - – Non-blocking
  - – Hit-under-miss
- ❑ Pipeline parallism
  - – ALU/MAC, LSU
  - – LS instruction won't stall the pipeline
  - – Out-of-order completion

Comparison

Feature	ARM9E™	ARM10E™	Intel® XScale™	ARM11TM
Architecture	ARMv5TE(J)	ARMv5TE(J)	ARMv5TE	ARMv6
Pipeline Length	5	6	7	8
Java Decode	(ARM926EJ)	(ARM1026EJ)	No	Yes
V6 SIMD Instructions	No	No	No	Yes
MIA Instructions	No	No	Yes	Available as coprocessor
Branch Prediction	No	Static	Dynamic	Dynamic
Independent Load-Store Unit	No	Yes	Yes	Yes
Instruction Issue	Scalar, in-order	Scalar, in-order	Scalar, in-order	Scalar, in-order
Concurrency	None	ALU/MAC, LSU	ALU, MAC, LSU	ALU/MAC, LSU
Out-of-order completion	No	Yes	Yes	Yes
Target Implementation	Synthesizable	Synthesizable	Custom chip	Synthesizable and Hard macro

Introduction to Several ARM processors
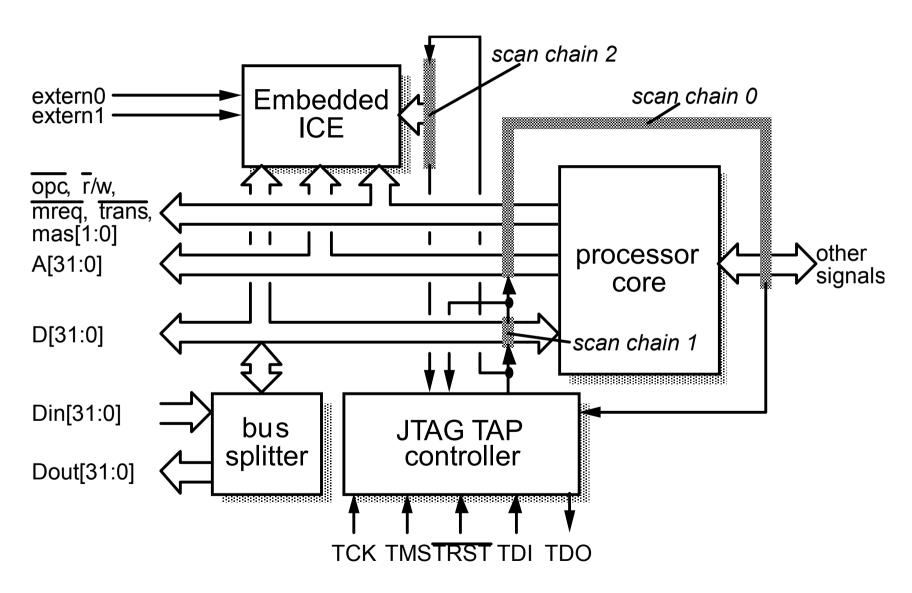
ARM7TDMI Processor Core

❑ Current low-end ARM core for applications like digital mobile phones

❑ TDMI

– **T**: Thumb, 16-bit compressed instruction set

– **D**: on-chip Debug support, enabling the processor to halt in response to a debug request

– **M**: enhanced Multiplier, yield a full 64-bit result, high performance

– **I**: Embedded ICE hardware
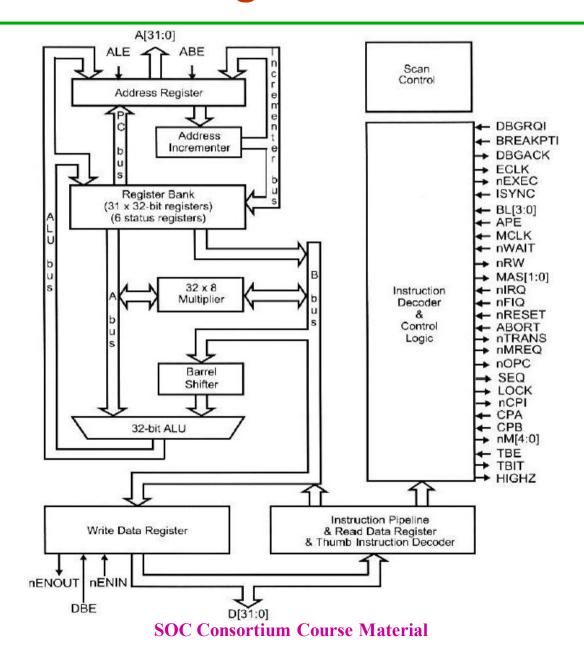
❑ Von Neumann architecture

❑ 3-stage pipeline, CPI ~ 1.9

ARM7TDMI Block Diagram

extern0
extern1

Embedded ICE

scan chain 2

scan chain 0

$\overline{opc}$, $\overline{r/w}$,
$\overline{mreq}$, $\overline{trans}$,
mas[1:0]

A[31:0]

processor core

other signals

D[31:0]

scan chain 1

Din[31:0]

bus splitter

JTAG TAP controller

Dout[31:0]

TCK TMS $\overline{TRST}$ TDI TDO
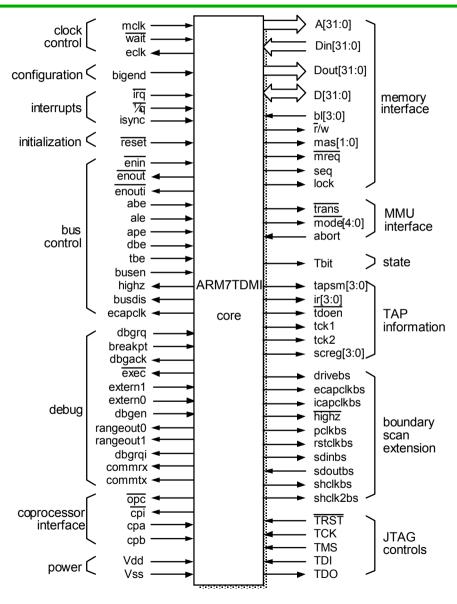
ARM7TDMI Core Diagram

ARM7TDMI Interface Signals (1/4)

ARM7TDMI Interface Signals (2/4)

❑ Clock control

- All state change within the processor are controlled by *mclk*, the memory clock
- Internal clock = mclk AND \wait
- eclk clock output reflects the clock used by the core

❑ Memory interface

- 32-bit address A[31:0], bidirectional data bus D[31:0], separate data out Dout[31:0], data in Din[31:0]
- \mreq indicates that the memory address will be sequential to that used in the previous cycle

mreq	seq	Cycle	Use
0	0	N	Non-sequential memory access
0	1	S	Sequential memory access
1	0	I	Internal cycle – bus and memory inactive
1	1	C	Coprocessor register transfer – memory inactive

ARM7TDMI Interface Signals (3/4)

- Lock indicates that the processor should keep the bus to ensure the atomicity of the read and write phase of a SWAP instruction
- \r/w, read or write
- mas[1:0], encode memory access size – byte, half–word or word
- bl[3:0], externally controlled enables on latches on each of the 4 bytes on the data input bus

❑ MMU interface

- \trans (translation control), 0: user mode, 1: privileged mode
- \mode[4:0], bottom 5 bits of the CPSR (inverted)
- Abort, disallow access

❑ State

- T bit, whether the processor is currently executing ARM or Thumb instructions

❑ Configuration

- Bigend, big-endian or little-endian

ARM7TDMI Interface Signals (4/4)

❑ Interrupt

– \fiq, fast interrupt request, higher priority

– \irq, normal interrupt request

– isync, allow the interrupt synchronizer to be passed

❑ Initialization

– \reset, starts the processor from a known state, executing from address 00000000_{16}
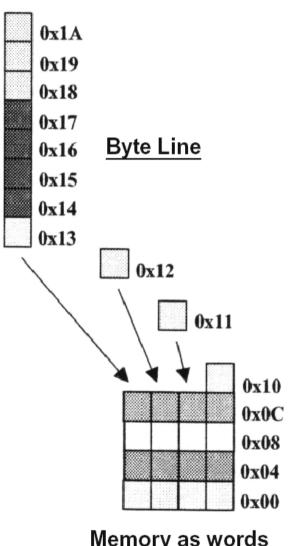
❑ ARM7TDMI characteristics

Process	0.35 um	Transistors	74,209	MIPS	60
Metal layers	3	Core area	2.1 mm^2	Power	87 mW
Vdd	3.3 V	Clock	0 to 66 MHz	MIPS/W	690

Memory Access

❑ The ARM7 is a Von Neumann, load/store architecture, i.e.,

  – Only 32 bit data bus for both instr. and data.
  – Only the load/store instr. (and SWP) access memory.

❑ Memory is addressed as a 32 bit address space

❑ Data type can be 8 bit bytes, 16 bit half-words or 32 bit words, and may be seen as a byte line folded into 4-byte words

❑ Words must be aligned to 4 byte boundaries, and half-words to 2 byte boundaries.

❑ Always ensure that memory controller supports all three access sizes

0x1A
0x19
0x18
0x17
0x16
0x15
0x14
0x13

Byte Line

0x12

0x11

0x10
0x0C
0x08
0x04
0x00

Memory as words
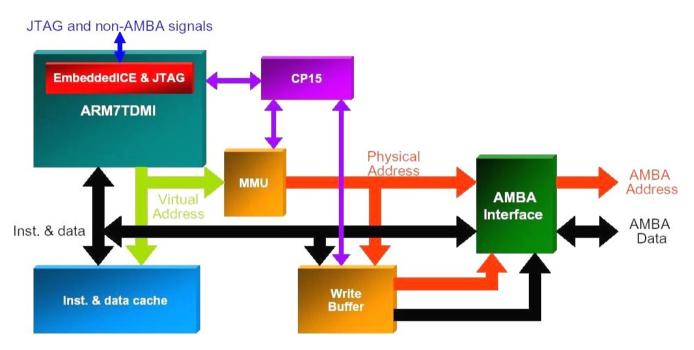
ARM Memory Interface

- Sequential (S cycle)
  - (nMREQ, SEQ) = (0, 1)
  - The ARM core requests a transfer to or from an address which is either the same, or one word or one-half-word greater than the preceding address.
- Non-sequential (N cycle)
  - (nMREQ, SEQ) = (0, 0)
  - The ARM core requests a transfer to or from an address which is unrelated to the address used in the preceding address.
- Internal (I cycle)
  - (nMREQ, SEQ) = (1, 0)
  - The ARM core does not require a transfer, as it performing an internal function, and no useful prefetching can be performed at the same time
- Coprocessor register transfer (C cycle)
  - (nMREQ, SEQ) = (1, 1)
  - The ARM core wished to use the data bus to communicate with a coprocessor, but does no require any action by the memory system.

Cached ARM7TDMI Macrocells

❑ ARM710T
- 8K unified write through cache
- Full memory management unit supporting virtual memory
- Write buffer

❑ ARM720T
- As ARM 710T but with WinCE support

❑ ARM 740T
- 8K unified write through cache
- Memory protection unit
- Write buffer
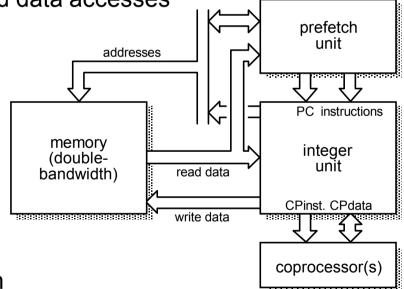
ARM8

❑ Higher performance than ARM7
  – By increasing the clock rate
  – By reducing the CPI
    • Higher memory bandwidth, 64-bit wide memory
    • Separate memories for instruction and data accesses

❑ ARM8 ⟶ ARM9TDMI
         ⟶ ARM10TDMI

❑ Core Organization
  – The prefetch unit is responsible for fetching instructions from memory and buffering them (exploiting the double bandwidth memory)
  – It is also responsible for branch prediction and use static prediction based on the branch prediction (backward: predicted 'taken'; forward: predicted 'not taken')

Pipeline Organization

❑ 5-stage, prefetch unit occupies the 1st stage, integer unit occupies the remainder

(1) Instruction prefetch   ⟶   **Prefetch Unit**

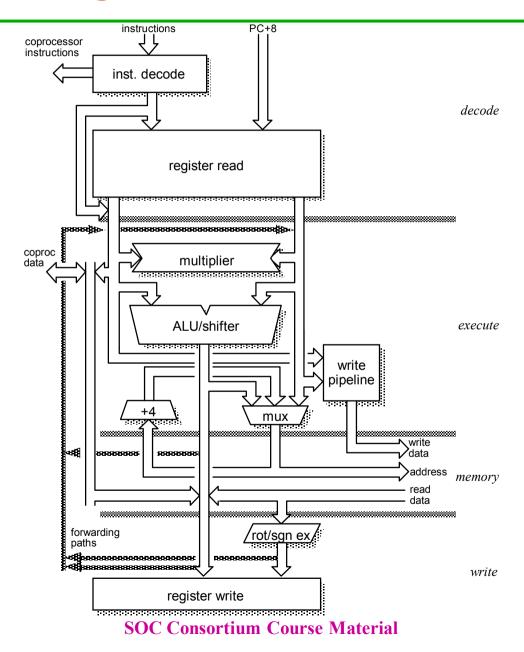(2) Instruction decode and register read

(3) Execute (shift and ALU)

(4) Data memory access   **Integer Unit**

(5) Write back results
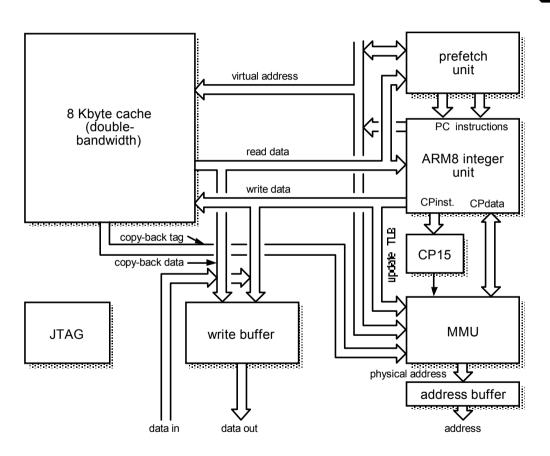
Integer Unit Organization

instructions PC+8

coprocessor instructions

inst. decode

decode

register read

coproc data

multiplier

ALU/shifter

execute

write pipeline

+4   mux

write data

address   *memory*

read data

forwarding paths

rot/sgn ex

write

register write

ARM8 Macrocell

❑ ARM810
   – 8Kbyte unified instruction and data cache
   – Copy-back
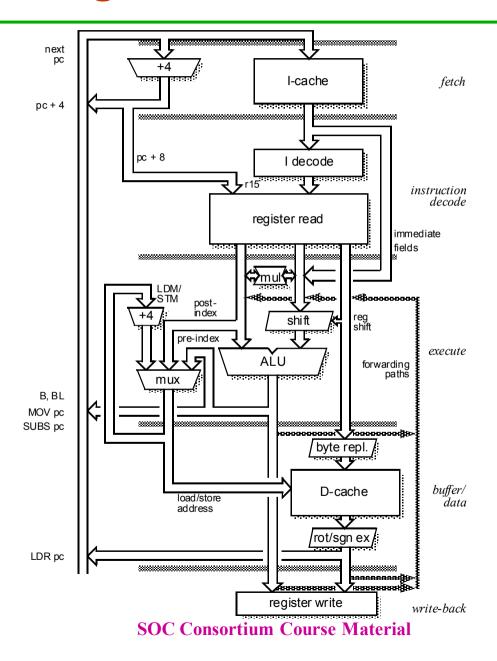   – Double-bandwidth
   – MMU
   – Coprocessor
   – Write buffer

Diagram labels:
- 8 Kbyte cache (double-bandwidth)
- virtual address
- read data
- write data
- copy-back tag
- copy-back data
- JTAG
- write buffer
- data in
- data out
- prefetch unit
- PC instructions
- ARM8 integer unit
- CPinst. CPdata
- update TLB
- CP15
- MMU
- physical address
- address buffer
- address

ARM9TDMI

❑ Harvard architecture

– Increases available memory bandwidth

• Instruction memory interface

• Data memory interface

– Simultaneous accesses to instruction and data memory can be achieved

❑ 5-stage pipeline

❑ Changes implemented to
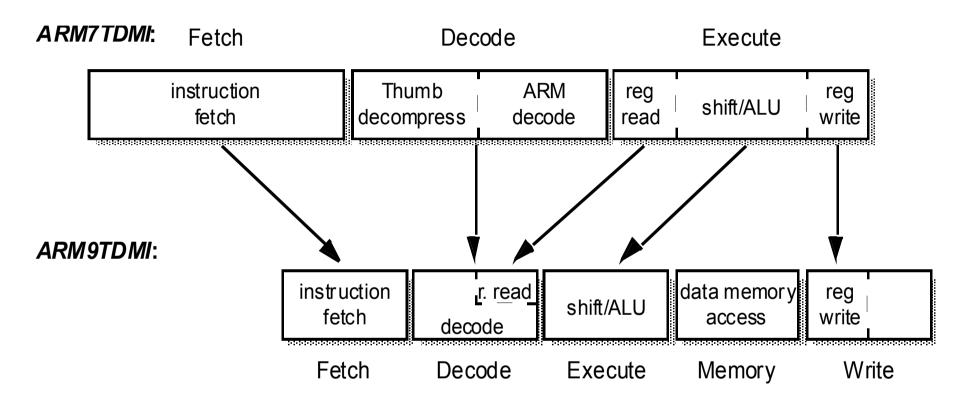
– Improve CPI to ~1.5

– Improve maximum clock frequency

ARM9TDMI Organization

ARM9TDMI Pipeline Operations (1/2)

ARM7TDMI:

Fetch   Decode   Execute

| instruction fetch | Thumb decompress | ARM decode | reg read | shift/ALU | reg write |

ARM9TDMI:

| instruction fetch | r. read / decode | shift/ALU | data memory access | reg write |

Fetch   Decode   Execute   Memory   Write

Not sufficient slack time to translate Thumb instructions into ARM instructions and then decode, instead the hardware decode both ARM and Thumb instructions directly

ARM9TDMI Pipeline Operations (2/2)

❑ Coprocessor support

– Coprocessors: floating-point, digital signal processing, special-purpose hardware accelerator

❑ On-chip debugger
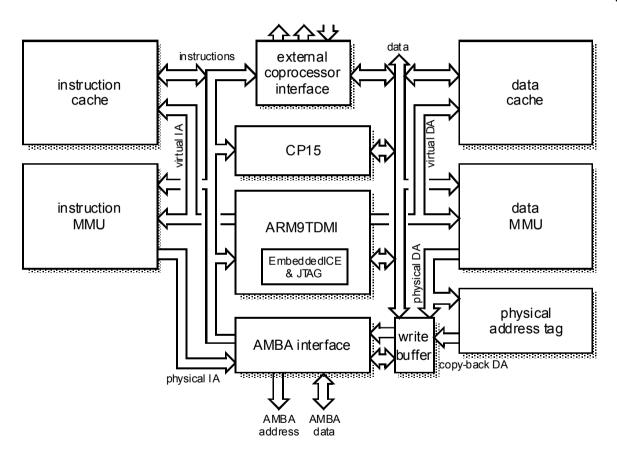
– Additional features compared to ARM7TDMI

  • Hardware single stepping

  • Breakpoint can be set on exceptions

❑ ARM9TDMI characteristics

Process	0.25 um	Transistors	110,000	MIPS	220
Metal layers	3	Core area	2.1 mm^2	Power	150 mW
Vdd	2.5 V	Clock	0 to 200 MHz	MIPS/W	1500

ARM9TDMI Macrocells (1/2)
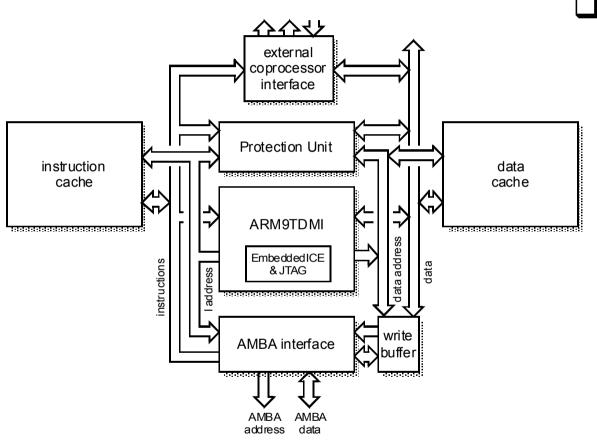
- ❑ ARM920T
  - 2 × 16K caches
  - Full memory management unit supporting virtual addressing and memory protection
  - Write buffer

ARM9TDMI Macrocells (2/2)

❑ ARM 940T
- 2 × 4K caches
- Memory protection Unit
- Write buffer

Diagram labels:
- external coprocessor interface
- instruction cache
- Protection Unit
- data cache
- ARM9TDMI
- Embedded ICE & JTAG
- instructions
- I address
- data address
- data
- AMBA interface
- write buffer
- AMBA address
- AMBA data

ARM9E-S Family Overview
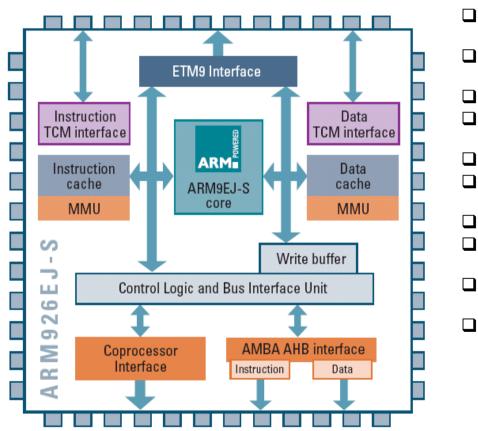
❑ ARM9E-S is based on an ARM9TDMI with the following extensions:

– Single cycle 32*6 multiplier implementation

– EmbeddedICE logic RT

– Improved ARM/Thumb interworking }

– New 32*16 and 16*16 multiply instructions } Architecture v5TE

– New count leading zero instruction

– New saturated math instructions

❑ ARM946E-S

– ARM9E-S core

– Instruction and data caches, selectable sizes

– Instruction and data RAMs, selectable sizes

– Protection unit

– AHB bus interface

ARM926EJ-S

- ❑ ARMv5TEJ architecture (ARMv5TEJ)
- ❑ 32-bit ARM instruction and 16-bit Thumb instruction set
- ❑ DSP instruction extensions and single cycle MAC
- ❑ ARM Jazelle technology
- ❑ MMU which supports operating systems including Symbian OS, Windows CE, Linux
- ❑ Flexible instruction and data cache sizes
- ❑ Instruction and data TCM interfaces with wait state support
- ❑ EmbeddedICE-RT logic for real-time debug
- ❑ Industry standard AMBA bus AHB interfaces
- ❑ ETM interface for Real-time trace capability with ETM9
- ❑ Optional MOVE Coprocessor delivers video encoding performance

ARM926EJ-S Performance Characteristics

	0.13um	0.18um
Area with cache (mm²)	3.2	8.3
Area w/o cache (mm²)	1.68	4.0
Frequency (MHz)	266	200-180
Typical mW/MHz with cache	0.45	1.40
Typical mW/MHz w/o cache	0.30	1.00

ARM10TDMI (1/2)
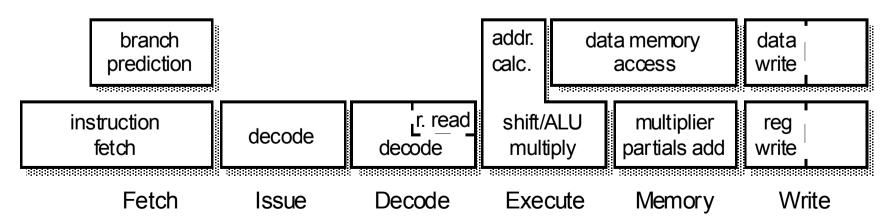
❑ Current high-end ARM processor core

❑ Performance on the same IC process

ARM10TDMI ⟵ ×2 ⟵ ARM9TDMI ⟵ ×2 ⟵ ARM7TDMI

❑ 300MHz, 0.25μm CMOS

❑ Increase clock rate

ARM10TDMI

branch prediction			addr. calc.	data memory access	data write
instruction fetch	decode	r. read decode	shift/ALU multiply	multiplier partials add	reg write
Fetch	Issue	Decode	Execute	Memory	Write

ARM10TDMI (2/2)

❑ Reduce CPI

- – Branch prediction

- – Non-blocking load and store execution

- – 64-bit data memory → transfer 2 registers in each cycle

ARM1020T Overview

❑ Architecture v5T
  – ARM1020E will be v5TE

❑ CPI ~ 1.3

❑ 6-stage pipeline

❑ Static branch prediction

❑ 32KB instruction and 32KB data caches
  – 'hit under miss' support

❑ 64 bits per cycle LDM/STM operations

❑ Embedded ICE Logic RT-II

❑ Support for new VFPv1 architecture

❑ ARM10200 test chip
  – ARM1020T
  – VFP10
  – SDRAM memory interface
  – PLL

ARM1176JZ(F)-S
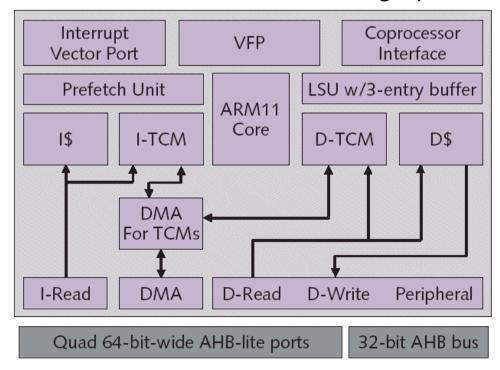
- ❑ Powerful ARMv6 instruction set architecture
  - – Thumb, Jazelle, DSP extensions
  - – SIMD (Single Instruction Multiple Data) media processing extensions deliver up to 2x performance for video processing
- ❑ Energy-saving power-down modes
  - – Reduce static leakage currents when processor is not in use
- ❑ High performance integer processor
  - – 8-stage integer pipeline delivers high clock frequency
  - – Separate load-store and arithmetic pipelines
  - – Branch Prediction and Return Stack
  - – Up to 660 Dhrystone 2.1 MIPS in 0.13µ process
- ❑ High performance memory system
  - – Supports 4-64k cache sizes
  - – Optional tightly coupled memories with DMA for multi-media applications
  - – **Multi-ported AMBA 2.0 AHB bus interface speeds instruction and data access**
  - – ARMv6 memory system architecture accelerates OS context-switch

ARM1176JZ(F)-S

❑ Vectored interrupt interface and low-interrupt-latency mode speeds interrupt response and real-time performance

❑ Optional Vector Floating Point coprocessor (ARM1136JF-S)

  – Powerful acceleration for embedded 3D-graphics

ARM1176JZ(F)-S Performance Characteristics

	0.13um
Area with cache (mm²)	5.55
Area w/o cache (mm²)	2.85
Frequency (MHz)	333-550
Typical mW/MHz with cache	0.8
Typical mW/MHz w/o cache	0.6

ARM11 MPCore
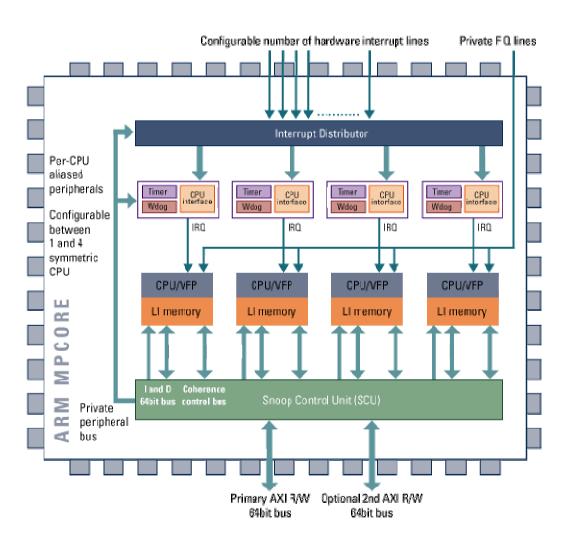
❑ Highly configurable
- – **Flexibility of total available performance from implementations using between 1 and 4 processors.**
- – Sizing of both data and instruction cache between 16K and 64K bytes across each processor.
- – Either dual or single 64-bit AMBA 3 AXI system bus connection allowing rapid and flexibility during SoC design
- – Optional integrated vector floating point (VFP) unit
- – Sizing on the number of hardware interrupts up to a total of 255 independent sources

ARM11 MPCore
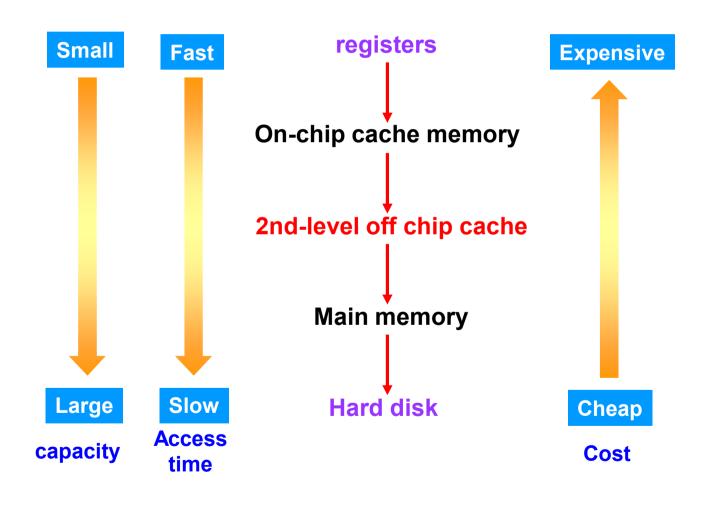
Memory Hierarchy

Memory Size and Speed

Small **Fast** registers **Expensive**

On-chip cache memory

2nd-level off chip cache

Main memory

Large **Slow** Hard disk **Cheap**

capacity **Access time** **Cost**

Caches (1/2)

❑ A cache memory is a small, very fast memory that retains copies of recently used memory values.

❑ It usually implemented on the same chip as the processor.

❑ Caches work because programs normally display the property of **locality**, which means that at any particular time they tend to execute the same instruction many times on the same areas of data.

❑ An access to an item which is in the cache is called a **hit**, and an access to an item which is not in the cache is a **miss**.
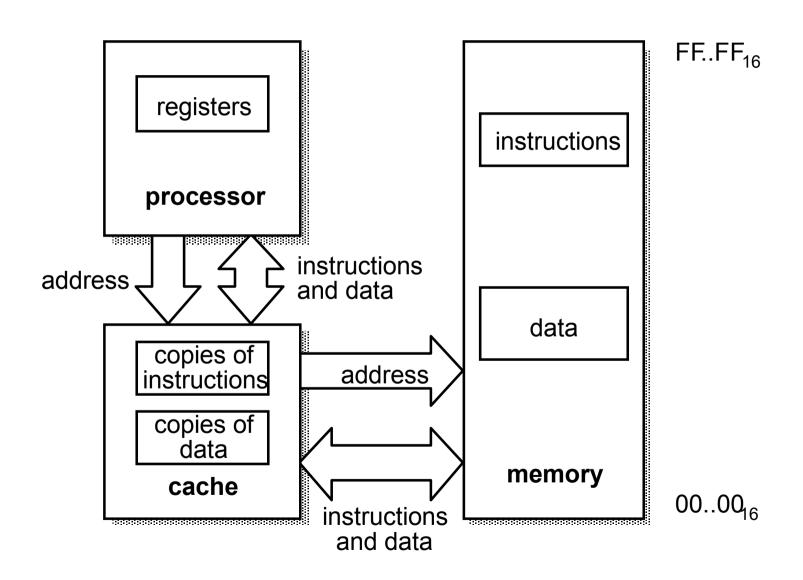
Caches (2/2)

❑ A processor can have one of the following two organizations:

– A unified cache

  • This is a single cache for both instructions and data

– Separate instruction and data caches

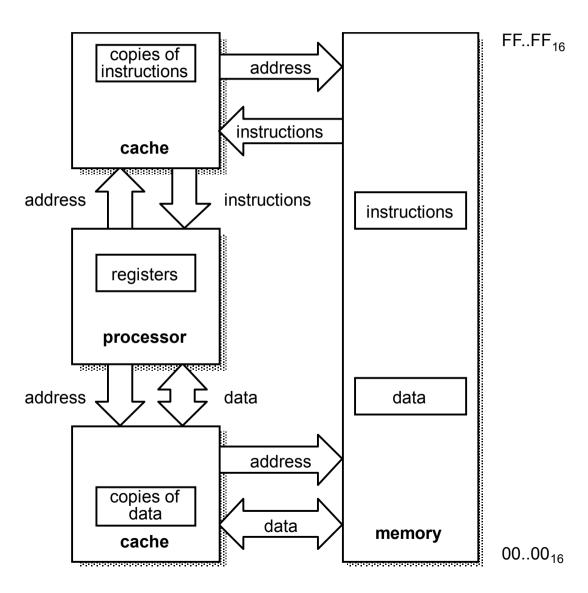 • This organization is sometimes called a **modified Harvard** architectures
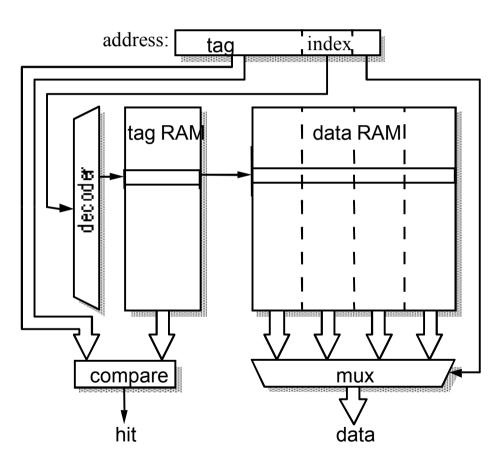
Unified Instruction and Data Cache

Separate Data and Instruction Caches

copies of instructions

address →
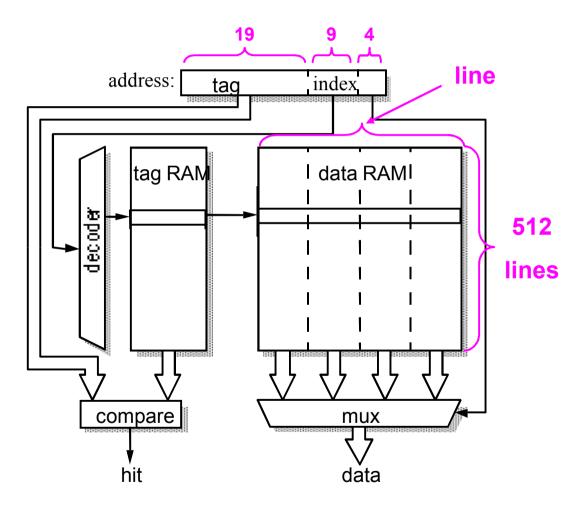
cache

← instructions

FF..FF$_{16}$

instructions

address ↑   instructions ↓

registers

processor

instructions

address ↓   data ↕

copies of data

cache

address →

data ↕

data

memory

00..00$_{16}$

The Direct-Mapped Cache

address: | tag | index |

tag RAM   data RAM

decoder

compare   mux

hit   data

- The index address bits are used to access the cache entry
- The top address bit are then compared with the stored tag
- If they are equal, the item is in the cache
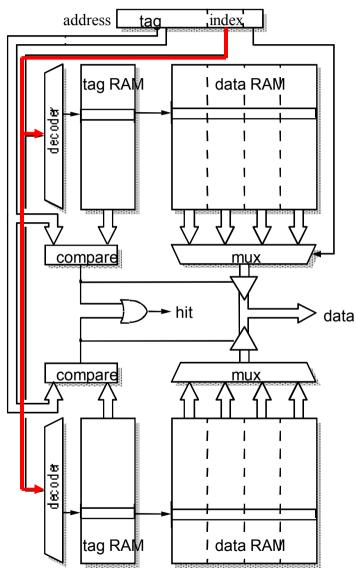- The lowest address bit can be used to access the desired item with in the line.

Example

19 **9** **4**

address: | tag | index |

line

512 lines

decoder

tag RAM

data RAM

compare

mux

hit

data

- ❏ The 8Kbytes of data in 16-byte lines. There would therefore be 512 lines

- ❏ A 32-bit address:
  - – 4 bits to address bytes within the line
  - – 9 bits to select the line
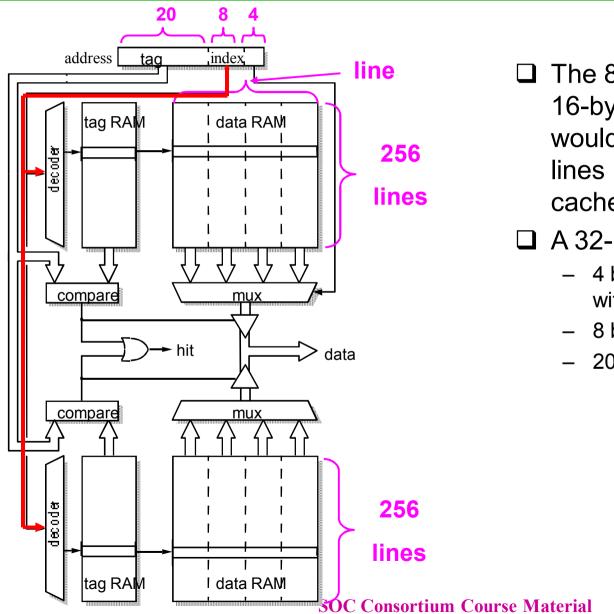 - – 19-bit tag

The Set-Associative Cache

- A 2-way set-associative cache
- This form of cache is effectively two direct-mapped caches operating in parallel.

Example

20   8   4

address   tag   index

line

tag RAM   data RAM

256 lines

decoder

compare   mux

hit   data
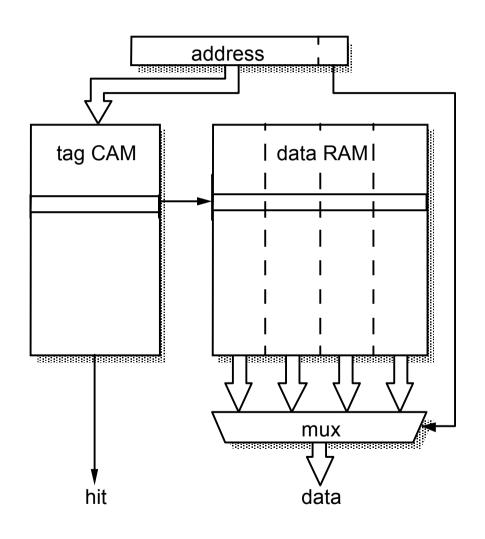
compare   mux

decoder

256 lines

tag RAM   data RAM

❑ The 8Kbytes of data in 16-byte lines. There would therefore be 256 lines in each half of the cache

❑ A 32-bit address:
  – 4 bits to address bytes within the line
  – 8 bits to select the line
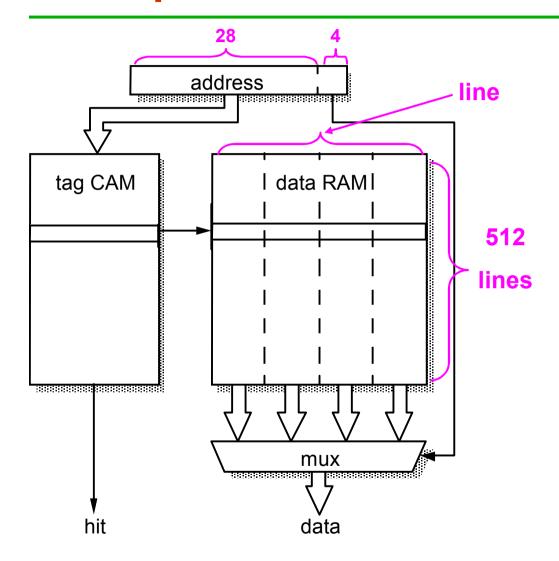 – 20-bit tag

Fully Associative Cache

- A **CAM** (Content Addressed Memory) cell is a RAM cell with an inbuilt comparator, so a CAM based tag store can perform a parallel search to locate an address in any location

- The address bit are compared with the stored tag

- If they are equal, the item is in the cache

- The lowest address bit can be used to access the desired item with in the line.

Example

28   **4**

address

line

tag CAM | data RAM |

512

lines

mux

hit   data

- ❑ The 8Kbytes of data in 16-byte lines. There would therefore be 512 lines
- ❑ A 32-bit address:
  - – 4 bits to address bytes within the line
  - – 28-bit tag

Write Strategies

❑ Write-through

– All write operations are passed to main memory

❑ Write-through with buffered write

– All write operations are still passed to main memory and the cache updated as appropriate, but instead of slowing the processor down to main memory speed the write address and data are stored in a **write buffer** which can accept the write information at high speed.
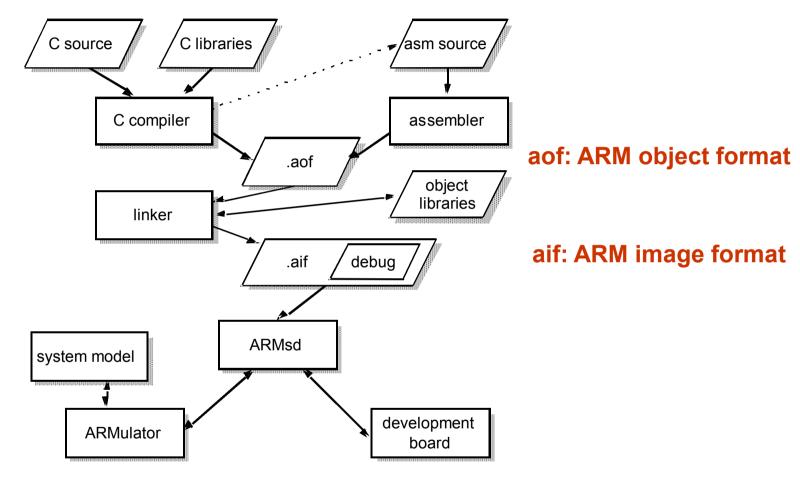
❑ Copy-back (write-back)

– No kept coherent with main memory

Software Development

ARM Tools

aof: ARM object format

aif: ARM image format

- ❑ ARM software development – *ADS*
- ❑ ARM system development – *ICE and trace*
- ❑ ARM-based SoC development – *modeling, tools, design flow*

ARM Development Suite (ADS), ARM Software Development Toolkit (SDT) (1/3)

❑ Develop and debug C/C++ or assembly language program

❑ *armcc* ARM C compiler

 armcpp   ARM C++ compiler

 tcc   Thumb C compiler

 tcpp   Thumb C++ compiler

 armasm   ARM and Thumb assembler

 armlink ARM linker

 armsd ARM and Thumb symbolic debugger

ARM Development Suite (ADS), ARM Software Development Toolkit (SDT) (2/3)

❑ *.aof*   ARM object format file

   .aif   ARM image format file

❑ The .aif file can be built to include the debug tables

- ARM symbolic debugger, ARMsd

❑ ARMsd can load, run and debug programs either on hardware such as the ARM development board or using the software emulation of the ARM

❑ AXD (ARM eXtended Debugger)

- ARM debugger for Windows and Unix with graphics user interface
- Debug C, C++, and assembly language source

CodeWarrior IDE

- Project management tool for windows

ARM Development Suite (ADS), ARM Software Development Toolkit (SDT) (3/3)

❑ Utilities

armprof   ARM profiler

Flash downloader   download binary images to Flash memory on a development board

❑ Supporting software

– *ARMulator* ARM core simulator

- Provide instruction accurate simulation of ARM processors and enable ARM and Thumb executable programs to be run on non-native hardware
- Integrated with the ARM debugger

– *Angle*   ARM debug monitor

- Run on target development hardware and enable you to develop and debug applications on ARM-based hardware

ARM C Compiler

❑ Compiler is compliant with the ANSI standard for C

❑ Supported by the appropriate library of functions

❑ Use ARM Procedure Call Standard, APCS for all external functions

– For procedure entry and exit

❑ May produce assembly source output

– Can be inspected, hand optimized and then assembled sequentially

❑ Can also produce Thumb codes

Linker

❑ Take one or more object files and combine them

❑ Resolve symbolic references between the object files and extract the object modules from libraries

❑ Normally the linker includes debug tables in the output file

ARM Symbolic Debugger

❑ A front-end interface to debug program running either under emulator (on the ARMulator) or remotely on a ARM development board (via a serial line or through JTAG test interface)

❑ ARMsd allows an executable program to be loaded into the ARMulator or a development board and run. It allows the setting of

– Breakpoints, addresses in the code

– Watchpoints, memory address if accessed as data address

  • Cause exception to halt so that the processor state can be examined

ARM Emulator (1/2)

❑ ARMulator is a suite of programs that models the behavior of various ARM processor cores in software on a host system

❑ It operates at various levels of accuracy

  – Instruction accuracy

  – Cycle accuracy

  – Timing accuracy

    • Instruction count or number of cycles can be measured for a program

    • Performance analysis

❑ Timing accuracy model is used for cache, memory management unit analysis, and so on

ARM Emulator (2/2)

❑ ARMulator supports a C library to allow complete C programs to run on the simulated system

❑ To run software on ARMulator, through ARM symbolic debugger or ARM GUI debuggers, AXD

❑ It includes

– Processor core models which can emulate any ARM core

– A memory interface which allows the characteristics of the target memory system to be modeled

– A coprocessor interface that supports custom coprocessor models

– An OS interface that allows individual system calls to be handled

ARM Development Board

❑ A circuit board including an ARM core (e.g. ARM7TDMI), memory component, I/O and electrically programmable devices

❑ It can support both hardware and software development before the final application-specific hardware is available

Summary (1/2)

- **ARM7TDMI**
  - Von Neumann architecture
  - 3-stage pipeline
  - CPI ~ 1.9

- **ARM9TDMI, ARM9E-S**
  - Harvard architecture
  - 5-stage pipeline
  - CPI ~ 1.5

- **ARM10TDMI**
  - Harvard architecture
  - 6-stage pipeline
  - CPI ~ 1.3

Summary (2/2)

❑ Cache

 – Direct-mapped cache

 – Set-associative cache

 – Fully associative cache

❑ Software Development

 – CodeWarrior

 – AXD

References

[1] http://twins.ee.nctu.edu.tw/courses/ip_core_02/index.html

[2] http://video.ee.ntu.edu.tw/~dip/slide.html

[2] **ARM System-on-Chip Architecture** by S.Furber, Addison Wesley Longman: ISBN 0-201-67519-6.

[3] www.arm.com